

POSTGRES PIPELINE: ESPECIALLY EMPHASIZING EXECUTION

by Paul Jungwirth

Illuminated Computing

May 2023

PHASES

- Parsing
- Analysis
- Rewriting
- Planning & Optimizing
- Executing

FOR EXAMPLE

```
ALTER TABLE ADD PERIOD valid_at (valid_from, valid_til)
```

```
PRIMARY KEY/UNIQUE (id, valid_at WITHOUT OVERLAPS)
```

```
UPDATE/DELETE FROM t  
  FOR PORTION OF valid_at  
  FROM '2020-01-01' TO '2030-01-01'
```

```
FOREIGN KEY (id, PERIOD valid_at)  
  REFERENCES parent (id, PERIOD valid_at)
```

FOR EXAMPLE

```
ALTER TABLE ADD PERIOD valid_at (valid_from, valid_til)

PRIMARY KEY/UNIQUE (id, valid_at WITHOUT OVERLAPS)

UPDATE/DELETE FROM t
  FOR PORTION OF valid_at
  FROM '2020-01-01' TO '2030-01-01'

FOREIGN KEY (id, PERIOD valid_at)
  REFERENCES parent (id, PERIOD valid_at)
```

FOR EXAMPLE

```
ALTER TABLE ADD PERIOD valid_at (valid_from, valid_til)

PRIMARY KEY/UNIQUE (id, valid_at WITHOUT OVERLAPS)

UPDATE/DELETE FROM t
  FOR PORTION OF valid_at
  FROM '2020-01-01' TO '2030-01-01'

FOREIGN KEY (id, PERIOD valid_at)
  REFERENCES parent (id, PERIOD valid_at)
```

FOR EXAMPLE

```
ALTER TABLE ADD PERIOD valid_at (valid_from, valid_til)

PRIMARY KEY/UNIQUE (id, valid_at WITHOUT OVERLAPS)

UPDATE/DELETE FROM t
  FOR PORTION OF valid_at
  FROM '2020-01-01' TO '2030-01-01'

FOREIGN KEY (id, PERIOD valid_at)
  REFERENCES parent (id, PERIOD valid_at)
```

FOR EXAMPLE

```
ALTER TABLE ADD PERIOD valid_at (valid_from, valid_til)

PRIMARY KEY/UNIQUE (id, valid_at WITHOUT OVERLAPS)

UPDATE/DELETE FROM t
  FOR PORTION OF valid_at
  FROM '2020-01-01' TO '2030-01-01'

FOREIGN KEY (id, PERIOD valid_at)
  REFERENCES parent (id, PERIOD valid_at)
```


FOR EXAMPLE

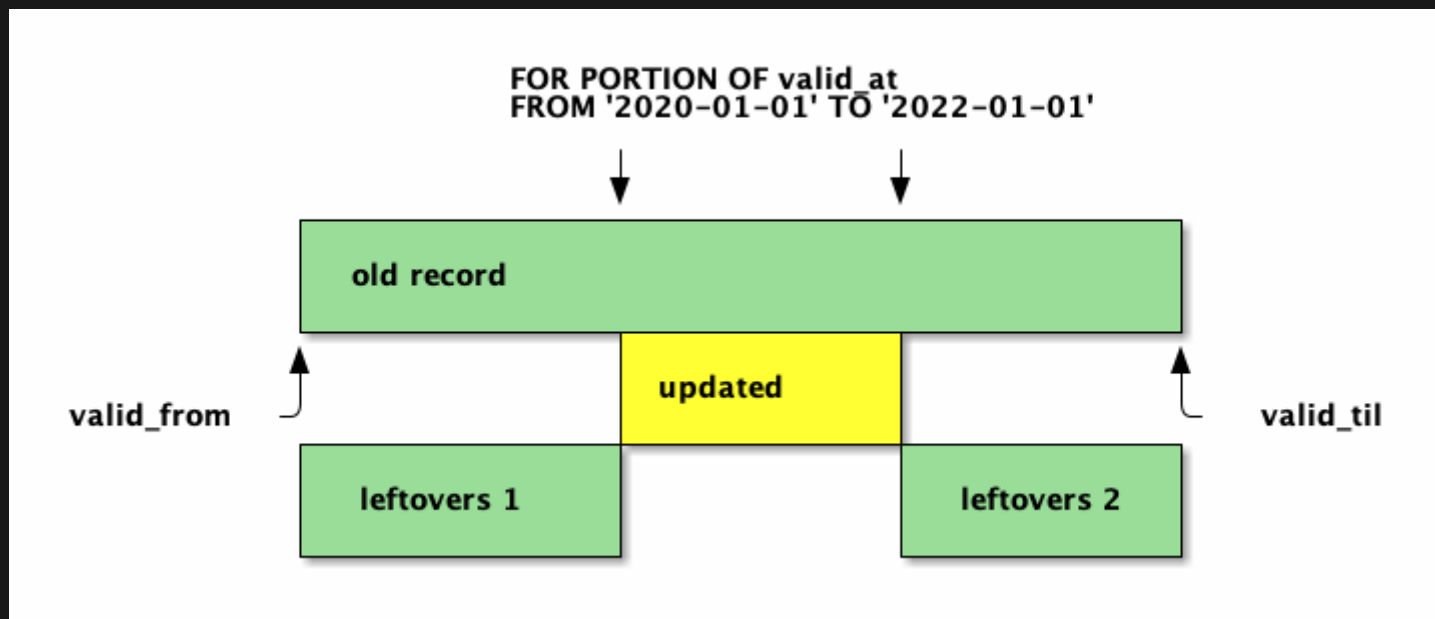
```
ALTER TABLE ADD PERIOD valid_at (valid_from, valid_til)

PRIMARY KEY/UNIQUE (id, valid_at WITHOUT OVERLAPS)

UPDATE/DELETE FROM t
  FOR PORTION OF valid_at
  FROM '2020-01-01' TO '2030-01-01'

FOREIGN KEY (id, PERIOD valid_at)
  REFERENCES parent (id, PERIOD valid_at)
```

FOR EXAMPLE



WHY

```
On Tue, Nov 14, 2017 at 9:43 AM Tom Lane <tgl@sss.pgh.pa.us> w
>
> Robert is correct that putting this into the parser
> is completely the wrong thing.
> If you do that, then for example views using the features
> will reverse-list in the rewritten form,
> which we Do Not Want,
> even if the rewritten form is completely valid SQL
> (is it?).
>
> . . .
>
>         regards, tom lane
```

TCOP/POSTGRES.C

```
exec_simple_query
  pg_parse_query
  pg_analyze_and_rewrite
    parse_analyze
    pg_rewrite_query
      QueryRewrite
        RewriteQuery
  pg_plan_queries
  PortalDefineQuery
  PortalStart
    ExecutorStart
      ExecInitModifyTable
  PortalRun
    FillPortalStore
    PortalRunSelect
      ExecutorRun
        ExecModifyTable
```

TCOP/POSTGRES.C

```
exec_simple_query
  pg_parse_query
  pg_analyze_and_rewrite
    parse_analyze
    pg_rewrite_query
      QueryRewrite
        RewriteQuery
  pg_plan_queries
  PortalDefineQuery
  PortalStart
    ExecutorStart
      ExecInitModifyTable
  PortalRun
    FillPortalStore
    PortalRunSelect
      ExecutorRun
        ExecModifyTable
```

TCOP/POSTGRES.C

```
exec_simple_query
  pg_parse_query
  pg_analyze_and_rewrite
    parse_analyze
    pg_rewrite_query
      QueryRewrite
        RewriteQuery
  pg_plan_queries
  PortalDefineQuery
  PortalStart
    ExecutorStart
      ExecInitModifyTable
  PortalRun
    FillPortalStore
    PortalRunSelect
      ExecutorRun
        ExecModifyTable
```

TCOP/POSTGRES.C

```
exec_simple_query
  pg_parse_query
  pg_analyze_and_rewrite
    parse_analyze
    pg_rewrite_query
      QueryRewrite
        RewriteQuery
  pg_plan_queries
  PortalDefineQuery
  PortalStart
    ExecutorStart
      ExecInitModifyTable
  PortalRun
    FillPortalStore
    PortalRunSelect
      ExecutorRun
        ExecModifyTable
```

TCOP/POSTGRES.C

```
exec_simple_query
  pg_parse_query
  pg_analyze_and_rewrite
    parse_analyze
    pg_rewrite_query
      QueryRewrite
      RewriteQuery
  pg_plan_queries
  PortalDefineQuery
  PortalStart
    ExecutorStart
      ExecInitModifyTable
  PortalRun
    FillPortalStore
    PortalRunSelect
      ExecutorRun
      ExecModifyTable
```


TCOP/POSTGRES.C

```
exec_simple_query
  pg_parse_query
  pg_analyze_and_rewrite
    parse_analyze
    pg_rewrite_query
      QueryRewrite
        RewriteQuery
  pg_plan_queries
  PortalDefineQuery
  PortalStart
    ExecutorStart
      ExecInitModifyTable
  PortalRun
    FillPortalStore
    PortalRunSelect
      ExecutorRun
        ExecModifyTable
```

TCOP/POSTGRES.C

```
exec_simple_query
  pg_parse_query
  pg_analyze_and_rewrite
    parse_analyze
    pg_rewrite_query
      QueryRewrite
        RewriteQuery
  pg_plan_queries
  PortalDefineQuery
  PortalStart
    ExecutorStart
      ExecInitModifyTable
  PortalRun
    FillPortalStore
    PortalRunSelect
      ExecutorRun
        ExecModifyTable
```

PARSING

```
/* src/backend/parser/gram.y */

for_portion_of_clause:
    FOR PORTION OF ColId FROM a_expr TO a_expr
    {
        ForPortionOfClause *n = makeNode(ForPortionOfClause);
        n->range_name = $4;
        n->range_name_location = @4;
        n->target_start = $6;
        n->target_end = $8;
        $$ = n;
    }
    | /*EMPTY*/           { $$ = NULL; }
;
```

PARSE NODES

```
/* src/include/nodes/parsenodes.h */
/*
 * ForPortionOfClause
 *     representation of FOR PORTION OF <period-name>
 *     FROM <t1> TO <t2>
 */
typedef struct ForPortionOfClause
{
    NodeTag      type;
    char         *range_name;
    int          range_name_location;
    Node         *target_start;
    Node         *target_end;
} ForPortionOfClause;
```

MEMORY CONTEXTS

```
TopMemoryContext  
PostmasterContext  
CacheMemoryContext  
MessageContext  
TopTransactionContext  
CurTransactionContext  
PortalContext  
ErrorContext  
and more!
```

ANALYSIS

```
static Query *
transformUpdateStmt(ParseState *pstate, UpdateStmt *stmt)
{
    Query          *qry = makeNode(Query);

    qry->resultRelation = setTargetTable(pstate, stmt->relation,
                                         stmt->relation->inh,
                                         true, ACL_UPDATE);

    if (stmt->forPortionOf)
        qry->forPortionOf = transformForPortionOfClause(
            pstate, qry->resultRelation,
            stmt->forPortionOf, true);

    transformFromClause(pstate, stmt->fromClause);
}
```

ANALYSIS

```
static Query *
transformUpdateStmt(ParseState *pstate, UpdateStmt *stmt)
{
    Query          *qry = makeNode(Query);

    qry->resultRelation = setTargetTable(pstate, stmt->relation,
                                         stmt->relation->inh,
                                         true, ACL_UPDATE);

    if (stmt->forPortionOf)
        qry->forPortionOf = transformForPortionOfClause(
            pstate, qry->resultRelation,
            stmt->forPortionOf, true);

    transformFromClause(pstate, stmt->fromClause);
}
```

ANALYSIS

```
static Query *
transformUpdateStmt(ParseState *pstate, UpdateStmt *stmt)
{
    Query          *qry = makeNode(Query);

    qry->resultRelation = setTargetTable(pstate, stmt->relation,
                                         stmt->relation->inh,
                                         true, ACL_UPDATE);

    if (stmt->forPortionOf)
        qry->forPortionOf = transformForPortionOfClause(
            pstate, qry->resultRelation,
            stmt->forPortionOf, true);

    transformFromClause(pstate, stmt->fromClause);
}
```


ANALYSIS

```
static Query *
transformUpdateStmt(ParseState *pstate, UpdateStmt *stmt)
{
    Query          *qry = makeNode(Query);

    qry->resultRelation = setTargetTable(pstate, stmt->relation,
                                         stmt->relation->inh,
                                         true, ACL_UPDATE);

    if (stmt->forPortionOf)
        qry->forPortionOf = transformForPortionOfClause(
            pstate, qry->resultRelation,
            stmt->forPortionOf, true);

    transformFromClause(pstate, stmt->fromClause);
}
```

ANALYSIS

```
static Query *
transformUpdateStmt(ParseState *pstate, UpdateStmt *stmt)
{
    Query          *qry = makeNode(Query);

    qry->resultRelation = setTargetTable(pstate, stmt->relation,
                                          stmt->relation->inh,
                                          true, ACL_UPDATE);

    if (stmt->forPortionOf)
        qry->forPortionOf = transformForPortionOfClause(
            pstate, qry->resultRelation,
            stmt->forPortionOf, true);

    transformFromClause(pstate, stmt->fromClause);
}
```

RANGETBENTRY

```
typedef struct RangeTblEntry
{
    NodeTag      type;

    RTEKind      rtekind;           /* see above */
    Oid           relid;            /* OID of the relation */
    char         relkind;          /* relation kind (see pg_class)
    int          rellockmode;      /* lock level that query requires
    Index         perminfoindex;
    Query        *subquery;        /* the sub-query */
    /* . . . */
};
```

SYSCACHE

```
HeapTuple perTuple = SearchSysCache2(PERIODNAME,  
                                     ObjectIdGetDatum(relid),  
                                     PointerGetDatum(range_name),  
                                     0, 0, 0, 0);  
  
if (HeapTupleIsValid(perTuple))  
{  
    Form_pg_period per = (Form_pg_period) GETSTRUCT(perTuple);  
    Oid rngtypeid      = per->perrngtype;  
    int start_attno    = per->perstart;  
    int end_attno      = per->perend;  
    Type rngtype       = typeidType(per->perrngtype);  
    char *range_type_name = typeTypeName(rngtype);  
    . . .  
    ReleaseSysCache(rngtype);  
    ReleaseSysCache(perTuple);  
}
```

SYSCACHE

```
HeapTuple perTuple = SearchSysCache2(PERIODNAME,  
                                     ObjectIdGetDatum(relid),  
                                     PointerGetDatum(range_name),  
                                     0, 0, 0, 0);  
  
if (HeapTupleIsValid(perTuple))  
{  
    Form_pg_period per = (Form_pg_period) GETSTRUCT(perTuple);  
    Oid rngtypeid      = per->perrngtype;  
    int start_attno    = per->perstart;  
    int end_attno      = per->perend;  
    Type rngtype       = typeidType(per->perrngtype);  
    char *range_type_name = typeTypeName(rngtype);  
    . . .  
    ReleaseSysCache(rngtype);  
    ReleaseSysCache(perTuple);  
}
```

SYSCACHE

```
HeapTuple perTuple = SearchSysCache2(PERIODNAME,  
                                     ObjectIdGetDatum(releid),  
                                     PointerGetDatum(range_name),  
                                     0, 0, 0, 0);  
  
if (HeapTupleIsValid(perTuple))  
{  
    Form_pg_period per = (Form_pg_period) GETSTRUCT(perTuple);  
    Oid rngtypeid      = per->perrngtype;  
    int start_attno    = per->perstart;  
    int end_attno      = per->perend;  
    Type rngtype       = typeidType(per->perrngtype);  
    char *range_type_name = typeTypeName(rngtype);  
    . . .  
    ReleaseSysCache(rngtype);  
    ReleaseSysCache(perTuple);  
}
```

SYSCACHE

```
HeapTuple perTuple = SearchSysCache2(PERIODNAME,  
                                     ObjectIdGetDatum(releid),  
                                     PointerGetDatum(range_name),  
                                     0, 0, 0, 0);  
  
if (HeapTupleIsValid(perTuple))  
{  
    Form_pg_period per = (Form_pg_period) GETSTRUCT(perTuple);  
    Oid rngtypeid      = per->perrngtype;  
    int start_attno    = per->perstart;  
    int end_attno      = per->perend;  
    Type rngtype       = typeidType(per->perrngtype);  
    char *range_type_name = typeTypeName(rngtype);  
    . . .  
    ReleaseSysCache(rngtype);  
    ReleaseSysCache(perTuple);  
}
```

SYSCACHE

```
HeapTuple perTuple = SearchSysCache2(PERIODNAME,  
                                     ObjectIdGetDatum(relid),  
                                     PointerGetDatum(range_name),  
                                     0, 0, 0, 0);  
  
if (HeapTupleIsValid(perTuple))  
{  
    Form_pg_period per = (Form_pg_period) GETSTRUCT(perTuple);  
    Oid rngtypeid      = per->perrngtype;  
    int start_attno    = per->perstart;  
    int end_attno      = per->perend;  
    Type rngtype       = typeidType(per->perrngtype);  
    char *range_type_name = typeTypeName(rngtype);  
    . . .  
    ReleaseSysCache(rngtype);  
    ReleaseSysCache(perTuple);  
}
```


SYSCACHE

```
HeapTuple perTuple = SearchSysCache2(PERIODNAME,  
                                     ObjectIdGetDatum(relid),  
                                     PointerGetDatum(range_name),  
                                     0, 0, 0, 0);  
  
if (HeapTupleIsValid(perTuple))  
{  
    Form_pg_period per = (Form_pg_period) GETSTRUCT(perTuple);  
    Oid rngtypeid      = per->perrngtype;  
    int start_attno    = per->perstart;  
    int end_attno      = per->perend;  
    Type rngtype       = typeidType(per->perrngtype);  
    char *range_type_name = typeTypeName(rngtype);  
    . . .  
    ReleaseSysCache(rngtype);  
    ReleaseSysCache(perTuple);  
}
```

SYSCACHE

```
HeapTuple perTuple = SearchSysCache2(PERIODNAME,  
                                     ObjectIdGetDatum(releid),  
                                     PointerGetDatum(range_name),  
                                     0, 0, 0, 0);  
  
if (HeapTupleIsValid(perTuple))  
{  
    Form_pg_period per = (Form_pg_period) GETSTRUCT(perTuple);  
    Oid rngtypeid      = per->perrngtype;  
    int start_attno    = per->perstart;  
    int end_attno      = per->perend;  
    Type rngtype       = typeidType(per->perrngtype);  
    char *range_type_name = typeTypeName(rngtype);  
    . . .  
    ReleaseSysCache(rngtype);  
    ReleaseSysCache(perTuple);  
}
```

SYSCACHE

```
HeapTuple perTuple = SearchSysCache2(PERIODNAME,  
                                     ObjectIdGetDatum(releid),  
                                     PointerGetDatum(range_name),  
                                     0, 0, 0, 0);  
  
if (HeapTupleIsValid(perTuple))  
{  
    Form_pg_period per = (Form_pg_period) GETSTRUCT(perTuple);  
    Oid rngtypeid      = per->perrngtype;  
    int start_attno    = per->perstart;  
    int end_attno      = per->perend;  
    Type rngtype       = typeidType(per->perrngtype);  
    char *range_type_name = typeTypeName(rngtype);  
    . . .  
    ReleaseSysCache(rngtype);  
    ReleaseSysCache(perTuple);  
}
```

SYSCACHE

```
HeapTuple perTuple = SearchSysCache2(PERIODNAME,  
                                     ObjectIdGetDatum(releid),  
                                     PointerGetDatum(range_name),  
                                     0, 0, 0, 0);  
  
if (HeapTupleIsValid(perTuple))  
{  
    Form_pg_period per = (Form_pg_period) GETSTRUCT(perTuple);  
    Oid rngtypeid      = per->perrngtype;  
    int start_attno    = per->perstart;  
    int end_attno      = per->perend;  
    Type rngtype       = typeidType(per->perrngtype);  
    char *range_type_name = typeTypeName(rngtype);  
    . . .  
    ReleaseSysCache(rngtype);  
    ReleaseSysCache(perTuple);  
}
```

LSYSCACHE

```
char *get_periodname(Oid periodid, bool missing_ok) {
    HeapTuple tp = SearchSysCache1(PERIODOID,
                                   ObjectIdGetDatum(periodid));

    if (HeapTupleIsValid(tp)) {
        Form_pg_period period_tup = (Form_pg_period) GETSTRUCT(tp);
        char *result = pstrdup(NameStr(period_tup->pername));
        ReleaseSysCache(tp);
        return result;
    }

    if (!missing_ok)
        elog(ERROR, "cache lookup failed for period %d", periodid);
    return NULL;
}
```

TYPCACHE

```
RangeType *r = DatumGetRangeTypeP(src->fp_targetRange);
TypeCacheEntry *typcache =
    lookup_type_cache(RangeTypeGetOid(r),
                      TYPECACHE_RANGE_INFO);
dst->fp_targetRange = datumCopy(src->fp_targetRange,
                                typcache->typbyval,
                                typcache->typlen);
```

TYPCACHE

```
RangeType *r = DatumGetRangeTypeP(src->fp_targetRange);
TypeCacheEntry *typcache =
    lookup_type_cache(RangeTypeGetOid(r),
                      TYPECACHE_RANGE_INFO);
dst->fp_targetRange = datumCopy(src->fp_targetRange,
                                typcache->typbyval,
                                typcache->typlen);
```

TYPCACHE

```
RangeType *r = DatumGetRangeTypeP(src->fp_targetRange);
TypeCacheEntry *typcache =
    lookup_type_cache(RangeTypeGetOid(r),
                      TYPECACHE_RANGE_INFO);
dst->fp_targetRange = datumCopy(src->fp_targetRange,
                                typcache->typbyval,
                                typcache->typlen);
```


REWRITING

- VIEWS
- RULES
- Query \rightarrow List (of Query)

REWRITING

```
foreach(lc, parsetree->forPortionOf->rangeSet)
{
    TargetEntry *tle = (TargetEntry *) lfirst(lc);
    TargetEntry *view_tle;

    if (tle->resjunk) continue;

    view_tle = get_tle_by_resno(view_targetlist, tle->resno);
    if (view_tle != NULL &&
        !view_tle->resjunk &&
        IsA(view_tle->expr, Var))
        tle->resno = ((Var *) view_tle->expr)->varattno;
    else
        elog(ERROR, "attribute number %d not found in view tar
}
```

REWRITING

```
foreach(lc, parsetree->forPortionOf->rangeSet)
{
    TargetEntry *tle = (TargetEntry *) lfirst(lc);
    TargetEntry *view_tle;

    if (tle->resjunk) continue;

    view_tle = get_tle_by_resno(view_targetlist, tle->resno);
    if (view_tle != NULL &&
        !view_tle->resjunk &&
        IsA(view_tle->expr, Var))
        tle->resno = ((Var *) view_tle->expr)->varattno;
    else
        elog(ERROR, "attribute number %d not found in view tar
}
```

REWRITING

```
foreach(lc, parsetree->iorportionor->rangeset)
{
    TargetEntry *tle = (TargetEntry *) lfirst(lc);
    TargetEntry *view_tle;

    if (tle->resjunk) continue;

    view_tle = get_tle_by_resno(view_targetlist, tle->resno);
    if (view_tle != NULL &&
        !view_tle->resjunk &&
        IsA(view_tle->expr, Var))
        tle->resno = ((Var *) view_tle->expr)->varattno;
    else
        elog(ERROR, "attribute number %d not found in view tar
}
```

PLANNING & OPTIMIZING

- `Query -> PlannedStmt`

EXECUTOR

```
PortalStart
  ExecutorStart
    CreateExecutorState
    InitPlan
      ExecInitNode
      ...
      ExecInitModifyTable
      ...
PortalRun
  PortalRunSelect
    ExecutorRun
      ExecProcNode
      ExecModifyTable
```

EXECUTOR

```
PortalStart
  ExecutorStart
    CreateExecutorState
    InitPlan
      ExecInitNode
      ...
      ExecInitModifyTable
      ...
PortalRun
  PortalRunSelect
    ExecutorRun
      ExecProcNode
      ExecModifyTable
```

EXECUTOR

```
PortalStart
  ExecutorStart
    CreateExecutorState
    InitPlan
      ExecInitNode
      ...
      ExecInitModifyTable
      ...
PortalRun
  PortalRunSelect
    ExecutorRun
      ExecProcNode
      ExecModifyTable
```


EXECUTOR

```
PortalStart
  ExecutorStart
    CreateExecutorState
    InitPlan
      ExecInitNode
      ...
      ExecInitModifyTable
      ...
PortalRun
  PortalRunSelect
    ExecutorRun
      ExecProcNode
      ExecModifyTable
```

EXECUTOR

```
PortalStart
  ExecutorStart
    CreateExecutorState
    InitPlan
      ExecInitNode
      ...
      ExecInitModifyTable
      ...
PortalRun
  PortalRunSelect
    ExecutorRun
      ExecProcNode
      ExecModifyTable
```

EXECUTOR

```
typedef struct PlanState
{
    pg_node_attr(abstract)
    NodeTag      type;
    Plan         *plan;
    EState       *state;
    ExecProcNodeMtd ExecProcNode;
    ...
}
```

EXECUTOR

```
typedef struct PlanState
{
    pg_node_attr(abstract)
    NodeTag      type;
    Plan         *plan;
    EState       *state;
    ExecProcNodeMtd ExecProcNode;
    ...
}
```

EXECUTOR

```
typedef struct PlanState
{
    pg_node_attr(abstract)
    NodeTag      type;
    Plan         *plan;
    EState       *state;
    ExecProcNodeMtd ExecProcNode;
    ...
}
```

EXECUTOR

```
typedef struct PlanState
{
    pg_node_attr(abstract)
    NodeTag      type;
    Plan         *plan;
    EState       *state;
    ExecProcNodeMtd ExecProcNode;
    ...
}
```

EXECUTOR

```
typedef struct PlanState
{
    pg_node_attr(abstract)
    NodeTag      type;
    Plan         *plan;
    EState       *state;
    ExecProcNodeMtd ExecProcNode;
    ...
}
```

EXECUTOR

```
ExprState *exprState = ExecPrepareExpr(  
    (Expr *) forPortionOf->targetRange, estate);  
Datum targetRange = ExecEvalExpr(exprState, econtext, &isNull)  
resultRelInfo->ri_forPortionOf->fp_targetRange = targetRange;
```


EXECUTOR

```
/* Initialize slot for the existing tuple */

resultRelInfo->ri_forPortionOf->fp_Existing =
    table_slot_create(resultRelInfo->ri_RelationDesc,
                      &mtstate->ps.state->es_tupleTable);

/* Create the tuple slots for INSERTing the leftovers */

resultRelInfo->ri_forPortionOf->fp_Leftover1 =
    ExecInitExtraTupleSlot(mtstate->ps.state, tupDesc,
                           &TTSOpsVirtual);
resultRelInfo->ri_forPortionOf->fp_Leftover2 =
    ExecInitExtraTupleSlot(mtstate->ps.state, tupDesc,
                           &TTSOpsVirtual);
```

EXECUTOR

```
/* Initialize slot for the existing tuple */

resultRelInfo->ri_forPortionOf->fp_Existing =
    table_slot_create(resultRelInfo->ri_RelationDesc,
                      &mtstate->ps.state->es_tupleTable);

/* Create the tuple slots for INSERTing the leftovers */

resultRelInfo->ri_forPortionOf->fp_Leftover1 =
    ExecInitExtraTupleSlot(mtstate->ps.state, tupDesc,
                           &TTSOpsVirtual);
resultRelInfo->ri_forPortionOf->fp_Leftover2 =
    ExecInitExtraTupleSlot(mtstate->ps.state, tupDesc,
                           &TTSOpsVirtual);
```

EXECUTOR

```
/* Initialize slot for the existing tuple */

resultRelInfo->ri_forPortionOf->fp_Existing =
    table_slot_create(resultRelInfo->ri_RelationDesc,
                      &mtstate->ps.state->es_tupleTable);

/* Create the tuple slots for INSERTing the leftovers */

resultRelInfo->ri_forPortionOf->fp_Leftover1 =
    ExecInitExtraTupleSlot(mtstate->ps.state, tupDesc,
                          &TTSOpsVirtual);
resultRelInfo->ri_forPortionOf->fp_Leftover2 =
    ExecInitExtraTupleSlot(mtstate->ps.state, tupDesc,
                          &TTSOpsVirtual);
```

EXECUTOR

```
/* Initialize slot for the existing tuple */

resultRelInfo->ri_forPortionOf->fp_Existing =
    table_slot_create(resultRelInfo->ri_RelationDesc,
                      &mtstate->ps.state->es_tupleTable);

/* Create the tuple slots for INSERTing the leftovers */

resultRelInfo->ri_forPortionOf->fp_Leftover1 =
    ExecInitExtraTupleSlot(mtstate->ps.state, tupDesc,
                           &TTSOpsVirtual);
resultRelInfo->ri_forPortionOf->fp_Leftover2 =
    ExecInitExtraTupleSlot(mtstate->ps.state, tupDesc,
                           &TTSOpsVirtual);
```

EXECUTOR

```
ExecutorRun
  ExecutePlan
    for (;;) {
      TupleTableSlot *slot = node->ExecProcNode(node);
      if (TupIsNull(slot))
        break;
    }
}
```

EXECUTOR

```
for (;;) {  
    context.planSlot = ExecProcNode(subplanstate);  
    if (TupIsNull(context.planSlot)) break;  
  
    switch (operation) {  
        case CMD_UPDATE:  
            slot = ExecUpdate(...);  
    }  
}
```

EXECUTOR

```
for (;;) {  
    context.planSlot = ExecProcNode(subplanstate);  
    if (TupIsNull(context.planSlot)) break;  
  
    switch (operation) {  
        case CMD_UPDATE:  
            slot = ExecUpdate(...);  
    }  
}
```

EXECUTOR

```
for (;;) {  
    context.planSlot = ExecProcNode(subplanstate);  
    if (TupIsNull(context.planSlot)) break;  
  
    switch (operation) {  
        case CMD_UPDATE:  
            slot = ExecUpdate(...);  
    }  
}
```


EXECUTOR

```
RangeType *oldRange = slot_getattr(oldtupleSlot,  
    forPortionOf->rangeVar->varattno, &isNull);  
RangeType *targetRange = DatumGetRangeTypeP(  
    resultRelInfo->ri_forPortionOf->fp_targetRange);  
  
range_leftover_internal(tpcache,  
    oldRangeType, targetRangeType,  
    &leftoverRangeType1, &leftoverRangeType2);
```

EXECUTOR

```
RangeType *oldRange = slot_getattr(oldtupleSlot,  
    forPortionOf->rangeVar->varattno, &isNull);  
RangeType *targetRange = DatumGetRangeTypeP(  
    resultRelInfo->ri_forPortionOf->fp_targetRange);  
  
range_leftover_internal(typpcache,  
    oldRangeType, targetRangeType,  
    &leftoverRangeType1, &leftoverRangeType2);
```

EXECUTOR

```
RangeType *oldRange = slot_getattr(oldtupleSlot,  
    forPortionOf->rangeVar->varattno, &isNull);  
RangeType *targetRange = DatumGetRangeTypeP(  
    resultRelInfo->ri_forPortionOf->fp_targetRange);  
  
range_leftover_internal(typcache,  
    oldRangeType, targetRangeType,  
    &leftoverRangeType1, &leftoverRangeType2);
```

EXECUTOR

```
RangeType *oldRange = slot_getattr(oldtupleSlot,  
    forPortionOf->rangeVar->varattno, &isNull);  
RangeType *targetRange = DatumGetRangeTypeP(  
    resultRelInfo->ri_forPortionOf->fp_targetRange);  
  
range_leftover_internal(typcache,  
    oldRangeType, targetRangeType,  
    &leftoverRangeType1, &leftoverRangeType2);
```

EXECUTOR

```
if (!RangeIsEmpty(leftoverRangeType1))
{
    HeapTuple oldtuple = ExecFetchSlotHeapTuple(
        oldtupleSlot, false, NULL);
    ExecForceStoreHeapTuple(oldtuple, leftoverTuple1, false);

    set_leftover_tuple_bounds(leftoverTuple1, forPortionOf,
        typcache, leftoverRangeType1);
    ExecMaterializeSlot(leftoverTuple1);

    ExecInsert(context, resultRelInfo, leftoverTuple1,
        node->canSetTag, NULL, NULL);
}
```

EXECUTOR

```
if (!RangeIsEmpty(leftoverRangeType1))
{
    HeapTuple oldtuple = ExecFetchSlotHeapTuple(
        oldtupleSlot, false, NULL);
    ExecForceStoreHeapTuple(oldtuple, leftoverTuple1, false);

    set_leftover_tuple_bounds(leftoverTuple1, forPortionOf,
                              typcache, leftoverRangeType1);
    ExecMaterializeSlot(leftoverTuple1);

    ExecInsert(context, resultRelInfo, leftoverTuple1,
               node->canSetTag, NULL, NULL);
}
```

EXECUTOR

```
if (!RangeIsEmpty(leftoverRangeType1))
{
    HeapTuple oldtuple = ExecFetchSlotHeapTuple(
        oldtupleSlot, false, NULL);
    ExecForceStoreHeapTuple(oldtuple, leftoverTuple1, false);

    set_leftover_tuple_bounds(leftoverTuple1, forPortionOf,
        typcache, leftoverRangeType1);
    ExecMaterializeSlot(leftoverTuple1);

    ExecInsert(context, resultRelInfo, leftoverTuple1,
        node->canSetTag, NULL, NULL);
}
```

EXECUTOR

```
if (!RangeIsEmpty(leftoverRangeType1))
{
    HeapTuple oldtuple = ExecFetchSlotHeapTuple(
        oldtupleSlot, false, NULL);
    ExecForceStoreHeapTuple(oldtuple, leftoverTuple1, false);

    set_leftover_tuple_bounds(leftoverTuple1, forPortionOf,
        typcache, leftoverRangeType1);
    ExecMaterializeSlot(leftoverTuple1);

    ExecInsert(context, resultRelInfo, leftoverTuple1,
        node->canSetTag, NULL, NULL);
}
```


EXECUTOR

```
if (!RangeIsEmpty(leftoverRangeType1))
{
    HeapTuple oldtuple = ExecFetchSlotHeapTuple(
        oldtupleSlot, false, NULL);
    ExecForceStoreHeapTuple(oldtuple, leftoverTuple1, false);

    set_leftover_tuple_bounds(leftoverTuple1, forPortionOf,
        typcache, leftoverRangeType1);
    ExecMaterializeSlot(leftoverTuple1);

    ExecInsert(context, resultRelInfo, leftoverTuple1,
        node->canSetTag, NULL, NULL);
}
```

EXECUTOR

```
if (!RangeIsEmpty(leftoverRangeType1))
{
    HeapTuple oldtuple = ExecFetchSlotHeapTuple(
        oldtupleSlot, false, NULL);
    ExecForceStoreHeapTuple(oldtuple, leftoverTuple1, false);

    set_leftover_tuple_bounds(leftoverTuple1, forPortionOf,
        typcache, leftoverRangeType1);
    ExecMaterializeSlot(leftoverTuple1);

    ExecInsert(context, resultRelInfo, leftoverTuple1,
        node->canSetTag, NULL, NULL);
}
```

TUPLE TABLE SLOTS

```
typedef struct TupleTableSlot
{
    NodeTag      type;
    uint16       tts_flags;          /* Boolean states */
    AttrNumber   tts_nvalid;         /* # of valid values in tts_va
    const TupleTableSlotOps *const tts_ops; /* implementation
    TupleDesc    tts_tupleDescriptor; /* slot's tuple descri
    Datum        *tts_values;        /* current per-attribute value
    bool         *tts_isnull;        /* current per-attribute isnul
    MemoryContext tts_mcxt;          /* slot itself is in this cont
    ItemPointerData tts_tid;         /* stored tuple's tid */
    Oid          tts_tableOid;       /* table oid of tuple */
} TupleTableSlot;
```

TUPLE TABLE SLOTS

```
typedef struct TupleTableSlot
{
    NodeTag      type;
    uint16       tts_flags;          /* Boolean states */
    AttrNumber    tts_nvalid;        /* # of valid values in tts_va
    const TupleTableSlotOps *const tts_ops; /* implementation
    TupleDesc     tts_tupleDescriptor; /* slot's tuple descri
    Datum         *tts_values;        /* current per-attribute value
    bool          *tts_isnull;        /* current per-attribute isnul
    MemoryContext tts_mcxt;           /* slot itself is in this cont
    ItemPointerData tts_tid;          /* stored tuple's tid */
    Oid           tts_tableOid;       /* table oid of tuple */
} TupleTableSlot;
```

TUPLE TABLE SLOTS

```
typedef struct TupleTableSlot
{
    NodeTag      type;
    uint16       tts_flags;          /* Boolean states */
    AttrNumber   tts_nvalid;         /* # of valid values in tts_va
    const TupleTableSlotOps *const tts_ops; /* implementation
    TupleDesc     tts_tupleDescriptor; /* slot's tuple descri
    Datum         *tts_values;        /* current per-attribute value
    bool          *tts_isnull;        /* current per-attribute isnul
    MemoryContext tts_mcxt;           /* slot itself is in this cont
    ItemPointerData tts_tid;          /* stored tuple's tid */
    Oid           tts_tableOid;       /* table oid of tuple */
} TupleTableSlot;
```

TUPLE TABLE SLOTS

```
typedef struct TupleTableSlot
{
    NodeTag      type;
    uint16       tts_flags;          /* Boolean states */
    AttrNumber   tts_nvalid;         /* # of valid values in tts_va
    const TupleTableSlotOps *const tts_ops; /* implementation
    TupleDesc     tts_tupleDescriptor; /* slot's tuple descri
    Datum         *tts_values;        /* current per-attribute value
    bool          *tts_isnull;         /* current per-attribute isnul
    MemoryContext tts_mcxt;            /* slot itself is in this cont
    ItemPointerData tts_tid;           /* stored tuple's tid */
    Oid           tts_tableOid;        /* table oid of tuple */
} TupleTableSlot;
```

TUPLE TABLE SLOTS

```
typedef struct TupleTableSlot
{
    NodeTag      type;
    uint16       tts_flags;          /* Boolean states */
    AttrNumber   tts_nvalid;         /* # of valid values in tts_va
    const TupleTableSlotOps *const tts_ops; /* implementation
    TupleDesc    tts_tupleDescriptor; /* slot's tuple descri
    Datum        *tts_values;        /* current per-attribute value
    bool         *tts_isnull;        /* current per-attribute isnul
    MemoryContext tts_mcxt;          /* slot itself is in this cont
    ItemPointerData tts_tid;         /* stored tuple's tid */
    Oid          tts_tableOid;       /* table oid of tuple */
} TupleTableSlot;
```

TUPLE TABLE SLOTS

```
typedef struct TupleTableSlot
{
    NodeTag      type;
    uint16       tts_flags;          /* Boolean states */
    AttrNumber   tts_nvalid;        /* # of valid values in tts_va
    const TupleTableSlotOps *const tts_ops; /* implementation
    TupleDesc     tts_tupleDescriptor; /* slot's tuple descri
    Datum         *tts_values;       /* current per-attribute value
    bool          *tts_isnull;       /* current per-attribute isnul
    MemoryContext tts_mcxt;          /* slot itself is in this cont
    ItemPointerData tts_tid;         /* stored tuple's tid */
    Oid           tts_tableOid;      /* table oid of tuple */
} TupleTableSlot;
```


TUPLE TABLE SLOTS

TTSOpsHeapTuple

```
typedef struct HeapTupleTableSlot
{
    pg_node_attr(abstract)
    TupleTableSlot base;
    HeapTuple    tuple;           /* physical tuple */
    uint32       off;            /* saved state for slot_deform
    HeapTupleData tupdata;       /* optional workspace for stor
} HeapTupleTableSlot;
```

TUPLE TABLE SLOTS

TTSOpsHeapTuple

```
typedef struct HeapTupleTableSlot
{
    pg_node_attr(abstract)
    TupleTableSlot base;
    HeapTuple    tuple;           /* physical tuple */
    uint32       off;            /* saved state for slot_deform
    HeapTupleData tupdata;       /* optional workspace for stor
} HeapTupleTableSlot;
```

TUPLE TABLE SLOTS

TTSOpsHeapTuple

```
typedef struct HeapTupleTableSlot
{
    pg_node_attr(abstract)
    TupleTableSlot base;
    HeapTuple    tuple;           /* physical tuple */
    uint32       off;            /* saved state for slot_deform
    HeapTupleData tupdata;       /* optional workspace for stor
} HeapTupleTableSlot;
```

TUPLE TABLE SLOTS

TTSOpsHeapTuple

```
RangeType *oldRange = slot_getattr(oldtupleSlot,  
    forPortionOf->rangeVar->varattno, &isNull);  
RangeType *targetRange = DatumGetRangeTypeP(  
    resultRelInfo->ri_forPortionOf->fp_targetRange);  
  
range_leftover_internal(typpcache,  
    oldRangeType, targetRangeType,  
    &leftoverRangeType1, &leftoverRangeType2);
```

TUPLE TABLE SLOTS

HeapTuple

```
HeapTuple perTuple = SearchSysCache2(PERIODNAME,
                                     ObjectIdGetDatum(relid),
                                     PointerGetDatum(range_name),
                                     0, 0, 0, 0);

if (HeapTupleIsValid(perTuple))
{
    Form_pg_period per = (Form_pg_period) GETSTRUCT(perTuple);
    Oid rngtypid      = per->perrngtype;
    int start_attno   = per->perstart;
    int end_attno     = per->perend;
    . . .
    ReleaseSysCache(perTuple);
}
```

TUPLE TABLE SLOTS

HeapTuple

```
HeapTuple perTuple = SearchSysCache2(PERIODNAME,  
                                     ObjectIdGetDatum(relid),  
                                     PointerGetDatum(range_name),  
                                     0,0,0,0);  
  
if (HeapTupleIsValid(perTuple))  
{  
    Form_pg_period per = (Form_pg_period) GETSTRUCT(perTuple);  
    Oid rngtypid      = per->perrngtype;  
    int start_attno   = per->perstart;  
    int end_attno     = per->perend;  
    . . .  
    ReleaseSysCache(perTuple);  
}
```

TUPLE TABLE SLOTS

TTSOpsBufferHeapTuple

```
typedef struct BufferHeapTupleTableSlot
{
    pg_node_attr(abstract)

    HeapTupleTableSlot base;

    Buffer          buffer;          /* tuple's buffer, or InvalidB
} BufferHeapTupleTableSlot;
```

TUPLE TABLE SLOTS

TTSlotsMinimalTuple

TUPLE TABLE SLOTS

TTSOpsVirtual

```
/* Initialize slot for the existing tuple */

resultRelInfo->ri_forPortionOf->fp_Existing =
    table_slot_create(resultRelInfo->ri_RelationDesc,
                      &mtstate->ps.state->es_tupleTable);

/* Create the tuple slots for INSERTing the leftovers */

resultRelInfo->ri_forPortionOf->fp_Leftover1 =
    ExecInitExtraTupleSlot(mtstate->ps.state, tupDesc,
                           &TTSOpsVirtual);
resultRelInfo->ri_forPortionOf->fp_Leftover2 =
    ExecInitExtraTupleSlot(mtstate->ps.state, tupDesc,
                           &TTSOpsVirtual);
```

TUPLE TABLE SLOTS

TTSOpsVirtual

```
/* Initialize slot for the existing tuple */

resultRelInfo->ri_forPortionOf->fp_Existing =
    table_slot_create(resultRelInfo->ri_RelationDesc,
                      &mtstate->ps.state->es_tupleTable);

/* Create the tuple slots for INSERTing the leftovers */

resultRelInfo->ri_forPortionOf->fp_Leftover1 =
    ExecInitExtraTupleSlot(mtstate->ps.state, tupDesc,
                           &TTSOpsVirtual);
resultRelInfo->ri_forPortionOf->fp_Leftover2 =
    ExecInitExtraTupleSlot(mtstate->ps.state, tupDesc,
                           &TTSOpsVirtual);
```

TUPLE TABLE SLOTS

TTSOpsVirtual

```
/* Initialize slot for the existing tuple */

resultRelInfo->ri_forPortionOf->fp_Existing =
    table_slot_create(resultRelInfo->ri_RelationDesc,
                      &mtstate->ps.state->es_tupleTable);

/* Create the tuple slots for INSERTing the leftovers */

resultRelInfo->ri_forPortionOf->fp_Leftover1 =
    ExecInitExtraTupleSlot(mtstate->ps.state, tupDesc,
                          &TTSOpsVirtual);
resultRelInfo->ri_forPortionOf->fp_Leftover2 =
    ExecInitExtraTupleSlot(mtstate->ps.state, tupDesc,
                          &TTSOpsVirtual);
```

TUPLE TABLE SLOTS

TTSOpsVirtual

```
/* Initialize slot for the existing tuple */

resultRelInfo->ri_forPortionOf->fp_Existing =
    table_slot_create(resultRelInfo->ri_RelationDesc,
                      &mtstate->ps.state->es_tupleTable);

/* Create the tuple slots for INSERTing the leftovers */

resultRelInfo->ri_forPortionOf->fp_Leftover1 =
    ExecInitExtraTupleSlot(mtstate->ps.state, tupDesc,
                          &TTSOpsVirtual);
resultRelInfo->ri_forPortionOf->fp_Leftover2 =
    ExecInitExtraTupleSlot(mtstate->ps.state, tupDesc,
                          &TTSOpsVirtual);
```

TUPLE TABLE SLOTS

TTSOpsVirtual

```
/* Initialize slot for the existing tuple */

resultRelInfo->ri_forPortionOf->fp_Existing =
    table_slot_create(resultRelInfo->ri_RelationDesc,
                      &mtstate->ps.state->es_tupleTable);

/* Create the tuple slots for INSERTing the leftovers */

resultRelInfo->ri_forPortionOf->fp_Leftover1 =
    ExecInitExtraTupleSlot(mtstate->ps.state, tupDesc,
                           &TTSOpsVirtual);
resultRelInfo->ri_forPortionOf->fp_Leftover2 =
    ExecInitExtraTupleSlot(mtstate->ps.state, tupDesc,
                           &TTSOpsVirtual);
```

TUPLE TABLE SLOTS

TTSOpsVirtual

```
HeapTuple oldtuple = ExecFetchSlotHeapTuple(  
    oldtupleSlot, false, NULL);  
ExecForceStoreHeapTuple(oldtuple, leftoverTuple1, false);  
  
set_leftover_tuple_bounds(leftoverTuple1, forPortionOf,  
                           typcache, leftoverRangeType1);  
ExecMaterializeSlot(leftoverTuple1);  
  
ExecInsert(context, resultRelInfo, leftoverTuple1,  
           node->canSetTag, NULL, NULL);
```

TUPLE TABLE SLOTS

TTSOpsVirtual

```
HeapTuple oldtuple = ExecFetchSlotHeapTuple(  
    oldtupleSlot, false, NULL);  
ExecForceStoreHeapTuple(oldtuple, leftoverTuple1, false);  
  
set_leftover_tuple_bounds(leftoverTuple1, forPortionOf,  
    typcache, leftoverRangeType1);  
ExecMaterializeSlot(leftoverTuple1);  
  
ExecInsert(context, resultRelInfo, leftoverTuple1,  
    node->canSetTag, NULL, NULL);
```

TUPLE TABLE SLOTS

TTSOpsVirtual

```
HeapTuple oldtuple = ExecFetchSlotHeapTuple(  
    oldtupleSlot, false, NULL);  
ExecForceStoreHeapTuple(oldtuple, leftoverTuple1, false);  
  
set_leftover_tuple_bounds(leftoverTuple1, forPortionOf,  
    typcache, leftoverRangeType1);  
ExecMaterializeSlot(leftoverTuple1);  
  
ExecInsert(context, resultRelInfo, leftoverTuple1,  
    node->canSetTag, NULL, NULL);
```


TUPLE TABLE SLOTS

TTSOpsVirtual

```
HeapTuple oldtuple = ExecFetchSlotHeapTuple(  
    oldtupleSlot, false, NULL);  
ExecForceStoreHeapTuple(oldtuple, leftoverTuple1, false);  
  
set_leftover_tuple_bounds(leftoverTuple1, forPortionOf,  
                           typcache, leftoverRangeType1);  
ExecMaterializeSlot(leftoverTuple1);  
  
ExecInsert(context, resultRelInfo, leftoverTuple1,  
           node->canSetTag, NULL, NULL);
```

TUPLE TABLE SLOTS

TTSOpsVirtual

```
HeapTuple oldtuple = ExecFetchSlotHeapTuple(  
    oldtupleSlot, false, NULL);  
ExecForceStoreHeapTuple(oldtuple, leftoverTuple1, false);  
  
set_leftover_tuple_bounds(leftoverTuple1, forPortionOf,  
    typcache, leftoverRangeType1);  
ExecMaterializeSlot(leftoverTuple1);  
  
ExecInsert(context, resultRelInfo, leftoverTuple1,  
    node->canSetTag, NULL, NULL);
```

TUPLE TABLE SLOTS

TTSOpsVirtual

```
HeapTuple oldtuple = ExecFetchSlotHeapTuple(  
    oldtupleSlot, false, NULL);  
ExecForceStoreHeapTuple(oldtuple, leftoverTuple1, false);  
  
set_leftover_tuple_bounds(leftoverTuple1, forPortionOf,  
    typcache, leftoverRangeType1);  
ExecMaterializeSlot(leftoverTuple1);  
  
ExecInsert(context, resultRelInfo, leftoverTuple1,  
    node->canSetTag, NULL, NULL);
```

TUPLE TABLE SLOTS

TTSOpsVirtual

```
HeapTuple oldtuple = ExecFetchSlotHeapTuple(  
    oldtupleSlot, false, NULL);  
ExecForceStoreHeapTuple(oldtuple, leftoverTuple1, false);  
  
set_leftover_tuple_bounds(leftoverTuple1, forPortionOf,  
                           typcache, leftoverRangeType1);  
ExecMaterializeSlot(leftoverTuple1);  
  
ExecInsert(context, resultRelInfo, leftoverTuple1,  
           node->canSetTag, NULL, NULL);
```

THANK YOU!

REFERENCES

1. Selena Deckelmann, *So, you want to a developer*, 2011.
https://wiki.postgresql.org/wiki/So,_you_want_to_be_a_developer%3F
2. Laetitia Avrot, *Demystifying Contributing to PostgreSQL*, 2018. <https://www.slideshare.net/LtitiaAvrot/demystifying-contributing-to-postgresql>
3. Neil Conway and Gavin Sherry, *Introduction to Hacking PostgreSQL*, 2007.
http://www.neilconway.org/talks/hacking/hack_slides.pdf and
https://www.cse.iitb.ac.in/infolab/Data/Courses/CS631/PostgreSQL-Resources/hacking_intro.pdf
4. Greg Smith, *Exposing PostgreSQL Internals with User-Defined Functions*, 2010.
https://www.pgcon.org/2010/schedule/attachments/142_HackingWithUDFs.pdf
5. Hironobu Suzuki, *The Internals of PostgreSQL*, 2012. <http://www.interdb.jp/pg/>
6. Egor Rogov, *Indexes in PostgreSQL*, 2019. <https://habr.com/ru/companies/postgrespro/articles/441962/>
7. Tom Lane, *Re: [HACKERS] [PROPOSAL] Temporal query processing with range types*, pgsql-hackers mailing list, 2017.
<https://www.postgresql.org/message-id/32265.1510681378@sss.pgh.pa.us>
8. Robert Haas, *Re: MERGE SQL statement for PG12*, pgsql-hackers mailing list, 2019. https://www.postgresql.org/message-id/CA%2BTgmoZj8fyJGAFxs%3D8Or9LeNyKe_xtoSN_zTeCSgoLrUye%3D9Q%40mail.gmail.com
9. Paul Jungwirth, <https://github.com/pjungwir/pgcon-2023-talk-exec-phase>

